



Abishek Subramanian

Associate Director
Databricks Technologies & Delivery

CONTACT

- Abishek_vnb@hotmail.com
- +91 98451 33922
- Bengaluru, India
- linkedin.com/in/abishek-subramanian
- github.com/learnwithabi
- learnwithabi.github.io

DATABRICKS & AI/ML

- Databricks Lakehouse
- Unity Catalog
- Delta Live Tables
- Apache Spark
- PySpark
- LakeFlow
- Genie AI/BI
- AgentBrick
- RAG / GenAI / LLM
- MLflow
- MLOps
- Model Serving
- Auto Loader
- Structured Streaming
- Databricks Apps
- Asset Bundles
- Delta Lake
- dbt
- ETL / ELT
- Databricks SQL

CLOUD PLATFORMS

- Azure Databricks
- ADLS Gen2
- Azure Event Hubs
- Azure Key Vault
- Azure DevOps
- ARM Templates
- AKS
- AWS Databricks
- Amazon S3 / EC2
- Amazon Bedrock
- AWS IAM
- GCP Databricks
- GCS / BigQuery
- Terraform

DEVOPS & PROGRAMMING

- Python
- SQL / Spark SQL
- Docker
- Kubernetes
- CI/CD
- Jenkins
- Ansible
- Bash
- Grafana
- ELK Stack
- Prometheus

MONITORING & AUTOMATION

- Grafana
- Prometheus
- ELK Stack
- SolarWinds
- KACE
- Automox
- Chocolatey

OPERATING SYSTEMS

- Ubuntu
- Red Hat / CentOS
- Oracle Linux
- Windows Server
- VMware ESXi
- XenServer

EDUCATION

B.Tech — Electronics & Communication Engineering
Pondicherry University
2004 — 2007 · Pondicherry, India

LANGUAGES

- English — Professional
- Tamil — Native

Abishek Subramanian

Page 1 of 2

Data & AI Solution Architect | Databricks Technologies & Delivery

CORE COMPETENCIES

Data Lakehouse Architecture · Unity Catalog & Governance · GenAI / RAG / LLM · Apache Spark / PySpark · ETL/ELT Pipeline Development · Delta Live Tables · MLflow & MLOps · Multi-Cloud Architecture (Azure · AWS · GCP) · Data Engineering · Python · CI/CD · Agile Delivery

PROFESSIONAL SUMMARY

Results-driven Data & AI leader with **18+ years** of enterprise experience architecting governed, AI-ready lakehouse platforms that deliver measurable business outcomes — at petabyte scale, across clouds, and at speed. Currently spearheading one of the region's most complex **Azure-to-GCP Databricks migrations** for a global retail enterprise (70+ TB, zero downtime). Previously at Databricks for **nearly 4 years**, partnering with Fortune 500 clients to design and operationalise high-impact data engineering and AI platforms. Deep expertise in the full Databricks stack — Unity Catalog, LakeFlow, Delta Live Tables, GenAI/RAG — and multi-cloud delivery across **Azure, AWS, and GCP**.

End-to-End Delivery Leadership

C-suite alignment to production deployment — strategy through execution

Databricks Lakehouse Architecture

Unity Catalog · LakeFlow · Delta Live Tables · GenAI at enterprise scale

Multi-Cloud Expertise

Azure · AWS · GCP — and hybrid / on-prem transitions at petabyte scale

CoE Buildout & Governance

Team scaling · security · cost optimisation · regulatory compliance

18+

Years Experience

5+

Years in Databricks

15+

Team Size Led

70+ TB

Data Migrated

20%+

Efficiency Gains

PROFESSIONAL EXPERIENCE

Associate Director — Databricks Technologies & Delivery

Feb 2026 – Present

Wavicle Data Solutions

Bengaluru, India

- Lead the **Databricks CoE (Center of Excellence)** — defining lakehouse architecture standards, governance frameworks, and delivery best practices across enterprise engagements
- Architecting a **petabyte-scale Azure-to-GCP Databricks lakehouse migration** for a global retail group — 70+ TB moved through a phased, fully automated, zero-downtime program
- Designed a **six-step automated migration framework**: asset classification → config mapping → JSON transform → dependency-ordered deployment → REST-API deployment → validation
- Built custom **Asset Migrator & Asset Validator** frameworks for governed, error-free cross-cloud migration of clusters, jobs, pipelines, secrets, libraries, and policies
- Implemented **Unity Catalog governance** on target platform — preserving access controls, data lineage, and audit logs throughout the migration lifecycle
- Lead a distributed offshore/onsite engineering team, owning solution architecture, technical mentoring, and senior stakeholder management

Senior Solution Engineer

May 2022 – Feb 2026

Databricks India Pvt. Ltd.

Bengaluru, India

- Established **Databricks Asset Bundle (DAB)** deployment best practices — reducing development cycles and infrastructure costs for enterprise customers
- Developed implementation guidelines for **GenAI using RAG architecture** (Genie AI/BI, AgentBrick), enabling product teams to adopt GenAI at enterprise scale
- Architected production-grade pipelines with **Delta Live Tables + dbt** — medallion architecture, data quality expectations, and automated Workflows orchestration
- Designed a **custom AI Bot Space** on Databricks Apps, integrating Unity Catalog and custom models into a secure, governed enterprise AI tool
- Built **ML lifecycle platforms** with MLflow — experiment tracking, model registry, REST-based Model Serving; resolved critical ML workload blockers for enterprise clients
- Optimised **platform costs via system tables** — DBU tracking, cluster utilisation analysis, and autoscaling policy enforcement across teams

CERTIFICATIONS

Databricks

Data Engineering Associate · Platform Administrator · Generative AI Fundamentals · Lakehouse Fundamentals

Microsoft Azure

Solutions Architect Expert · DevOps Engineer Expert · Administrator Associate · Identity & Access Administrator

AWS

Solutions Architect Professional · Solutions Architect Associate · AWS Community Builder

Oracle

VM 3.0 Implementation Specialist · Cloud Infrastructure Classic Associate Architect

Others

Red Hat Certified Engineer (RHCE) · VMware VCP-DCV · Citrix XenServer Administration · ITIL V3

EDUCATION

B.Tech — Electronics & Communication Engineering

Pondicherry University

2004 – 2007 · Pondicherry, India

LANGUAGES

English — Professional Tamil — Native

PROFILE HIGHLIGHTS

- 18+ years enterprise technology leadership across Databricks, cloud platforms, and large-scale data engineering
- Proven track record delivering 20%+ automation efficiency gains for global enterprises
- Built and led cross-functional teams of up to 15 engineers across offshore & onsite locations
- Deep expertise in regulated industries — security, compliance, and data governance at petabyte scale

SPEAKING & COMMUNITY

- Founder & Lead** — Databricks User Group India (India's largest; 6 cities: Bangalore, Hyderabad, Pune, Coimbatore, Kolkata, Chennai)
- Creator** — #LearnWithAbi (Databricks & GenAI content on LinkedIn, Medium, GitHub)
- Speaker — DATA+AI Summit Recap: *Databricks & Apache Iceberg*
- Speaker — 9th World ML Summit: *"Unlocking ML Ops Potential"*
- Speaker — ServerlessDays Bengaluru: *Databricks Serverless SQL*
- Speaker — TechXConf 2024, Asia's largest AI & cloud conference
- Session — Nasdaq AWS Community: *Bedrock + Databricks Model Serving*
- Led *"Mastering Generative AI"* — building & serving models with Databricks
- Workshops at PSG Tech Coimbatore, Adamas University Kolkata & MITS
- Nominated — Topmate Creator Awards 2024, Extraordinaire Category

PROFESSIONAL EXPERIENCE (CONTINUED)

Principal Software Engineer

Oracle India Pvt. Ltd.

Bengaluru, India

Aug 2021 – May 2022

- Onboarded Oracle Cloud Fusion (SaaS) enterprise customers — owning solution proposals, implementation, and cross-stakeholder coordination from pre-sales through go-live
- Defined cloud migration strategy from on-premises to **Oracle Cloud Infrastructure (OCI)** for SaaS and IaaS workloads across enterprise environments
- Drove escalations to swift resolution by leveraging deep technical capability to pre-empt delivery issues and coordinate effectively with engineering leads

System Integration Specialist Advisor

NTT Data

May 2018 – Aug 2021

- Designed scalable, secure, and cost-effective multi-cloud architectures across **Azure, AWS, and VMware**; led legacy-to-cloud-native modernisation projects
- Built IaC automation frameworks using **Terraform, Ansible, and Jenkins** — streamlining provisioning, deployment pipelines, and configuration management
- Established cloud governance, security, and compliance standards; collaborated with DevOps, security, and business stakeholders to align cloud strategy with organisational goals

EARLIER CAREER

Senior Software Engineer — Oracle India · Led Cloud Central team for SaaS Fusion post-provision assessments Jul 2015 – May 2018

Cloud Engineer — Citrix R&D India · XenApp, XenDesktop, NetScaler & StoreFront enterprise deployments Jun 2013 – Jul 2015

System Engineer — CSS Corp · Global Red Hat / XEN / VMware server administration Jun 2011 – Jun 2013

Software Engineer (Contract) — Wipro Technologies · UNIX batch ops, HPSM/HPOM automation, production support Sep 2010 – Jun 2011

KEY PROJECTS

Petabyte-Scale Azure-to-GCP Lakehouse Migration

70+ TB zero-downtime migration for a global retail enterprise with full Unity Catalog governance and six-step automated Asset Migrator & Validator framework.

Delta Live Tables + dbt Enterprise Pipeline

Medallion architecture pipelines with data quality expectations and automated Workflows orchestration for enterprise ELT at scale.

Amazon Bedrock + Databricks Model Serving

Integration between Bedrock foundation models and Databricks Model Serving endpoints, demonstrated at Nasdaq AWS Community.

Unity Catalog Governance Framework

Enterprise UC implementation — three-level namespace, fine-grained access controls, row/column-level security, data lineage, and audit logging across multi-workspace deployments.

DAB CI/CD Deployment Framework

Databricks Asset Bundle CI/CD best practices — automated workspace deployment, environment promotion, and version-controlled Databricks Jobs, pipelines, and clusters.

GenAI RAG AI Bot Space on Databricks Apps

Enterprise AI Bot Space using RAG, Unity Catalog-managed models, and MLflow for secure, governed GenAI workflows at scale.

Databricks Platform Cost Optimisation

System-tables observability — DBU tracking, cluster utilisation analysis, and autoscaling policies delivering measurable spend reduction.

Multi-Cloud IaC Automation Framework

Terraform, Ansible & Jenkins-based automation across Azure, AWS, VMware — provisioning, CI/CD pipelines, and cloud governance at NTT Data.

Structured Streaming Real-Time Pipeline

High-throughput event ingestion using Structured Streaming & Auto Loader — exactly-once semantics, schema evolution, and Delta Lake as sink for real-time analytics.

Databricks Security & Access Control

Enterprise-grade security implementation — cluster encryption, Unity Catalog row/column-level security, secret management, and compliance policy enforcement across workspaces.

RECOGNITION & COMMUNITY IMPACT

Founder & Lead

India's largest Databricks User Group — 6 cities, thousands of

Content Creator

#LearnWithAbi — Databricks & GenAI on LinkedIn, Medium,

AWS Community Builder

Active contributor to the AWS global ecosystem